

# Dynamic (run-time) partial reconfiguration of FPGAs

Andrei George-Andrei

Faculty of Electronics, Telecommunications and Information Technology

(Advanced Computing in Embedded Systems Master Program)

University POLITEHNICA of Bucharest

(Splaiul Independentei 313, Bucharest, Romania, 060042)

andrei.george@gmail.com

February 2022

**Abstract.** This paper will describe the state of the art of the area of partial reconfiguration computing. In the last years, the Field Programmable Gate Array has encountered an exponential growth of interest. From using the FPGAs only for Application Specific Integrated Circuits (ASIC) design and validation, now the trend is to use the FPGA directly in the final product as a standalone computing unit. The advantage of this practice is that the ASIC design, validation and manufacturing time and costs will be drastically reduced or even disappear, because the time- on-the-market of the product has the highest priority. State of the art FPGAs have evolved from simple gate arrays to complex systems with high speeds that competes with the classical microcontrollers (MCU), that include embedded memory, digital signal processing units (DSP), clock multipliers made of phase lock loops (PLL), tens of thousands of look up tables (LUT) and even microprocessors. These days, the FPGAs come along with different reconfiguration options. The most challenging of them is the partial run-time reconfiguration. This technique gives the system the versatility to update its hardware configuration on-the-fly, after the device has booted and started running. Also called Dynamic Partial Reconfiguration, it will provide the system with the ability to parallelize its resources and will gain a software-like ability to perform scheduled tasks, but this time by rearranging its hardware by the program demand. This will be accomplished without any performance reduction. Important areas of Dynamic Partial Reconfiguration interest can be the ones like runtime adaptive super-computing, fault tolerant systems, adaptive embedded software accelerators, self-reconfigurable and self-arrangeable systems.

## 1. Introduction

Traditionally, Application Specific Integrated Circuits (ASICs) are preferred for the final product. The prototyping and testing of the ASICs is accomplished by the help of Field Programmable Gate Array devices. The ASICs are the best choice for the mass production processes, because the cost of the production masks is being amortized by the large number of sold units. The last trend, on the other side, when speaking of low serial production, the Field Programmable Gate Arrays are preferred. This is because the development time is reduced, and the greatest part of the production price is reduced by eliminating the need of the production masks. These being said, the FPGAs are the best choice when talking about a low serial production and custom designs, because of their hardware versatility and ability to emulate almost every existing digital circuit. The Field Programmable Gate Arrays also come with a series of advantages, like high computing speeds, capable memory blocks, Digital Signal Processing units and even microprocessors. Also, the existence of the Intellectual Property (IP) Cores offered by the vendors are drastically reducing the time and costs of development of the final product [1] [9]. One major feature of the FPGAs that this paper will be discussing is the Partial Reconfiguration ability. This is the ability of the FPGA to change its hardware component structure at the run time. One advantage could be the parallelization of the tasks in a system. Let's consider a pure software system running on a Real Time Operating System (RTOS): all the hardware resources (Random Access Memory, processors, communications, peripherals, etc.) are already allocated and the configuration will remain the same for the whole operation time. Let's further assume that one of the subsystems will crash, resulting in the partial or total system failure. A dynamic partial reconfigurable FPGA-based system can be a better implementation of the earlier example. The main advantage of the run-time reconfigurable system is that while it is running it can easily change its hardware configuration, replacing pieces of subsystems or implementing them in different regions of the FPGA. Another advantage of this kind of system could be that in comparison with the RTOS tasks, the reconfigurable system can allocate or deallocate hardware subsystems in such way that only when they will be needed (for example as a periodic or sporadic task) they will be implemented in the circuit (they will be allocated when their functionality is needed and when they finish their job, they will be deallocated, erased from the hardware, and another circuit for another task can replace them). When the hardware resources are limited, or some low power requirements must be accomplished or even if the system is running in harsh (ex: radiation) [8] environment that can damage certain subsystems or subsystems parts, the Dynamic Partial Reconfiguration Systems can be the best choice. These reconfigurable can provide a high advantage in terms of power saving, task parallelization or even self-recovery when talking about fault tolerant systems [2].

## 2. Reconfiguration Techniques and Systems

In order to better understand the concept of Partial Reconfiguration (PR), the following terms must be discussed: granularity, reconfiguration rate, system coupling and types of reconfiguration techniques. System data transfer, latency, power consumption, throughput and overhead are some parameters influenced by the above properties [1] [2].

## 2.1. Reconfiguration Techniques


FPGAs are configured with a sequence of bits called bitstream. In the past, the FPGAs were configured with a single static bitstream as a full configured system. On the other hand, on these days the vendors offer some reconfiguration techniques that allow us to perform partial or full reconfiguration of the FPGA bitstream [1]:

- **Full Reconfiguration:** also known as total or global reconfiguration, this device programming method is a dynamic full reconfiguration that allows the configuration of the entire device without the need to reboot the external circuitry.

The main advantage is that this type of reconfiguration is supported by most of the FPGAs available on the market.

The disadvantage, although is that the bitstreams have large sizes (because the whole configuration is replaced). The consequences of the large bitstream sizes are that the reconfiguration time will be long (seconds), when reconfiguring, the power consumption will increase, and the storage demands will increase. The main disadvantage is that the whole system will be suspended while the reconfiguration is being done.

- **Partial Reconfiguration:** this type of reconfiguration allows to modify only on portion of the FPGA. A simple concept can be seen in Figure 1.



**Fig. 1.** Partial Reconfiguration Concept [4]

Furthermore, this type can be classified into two subcategories:

- **Non-disruptive (runtime or dynamic):** this type implies that the parts that are not being replaced will remain operational at the run-time. The advantage of this method is the reduced reconfiguration time ranging from milliseconds to a few seconds.
- **Disruptive (non-dynamic):** when this partial reconfiguration technique is applied, the rest of the system will be affected due to the clock dependencies that will hold the clock or input/output ports dependencies where the ports must be kept in high impedance. The advantage of this technique is reduced memory required for bitstreams and less power consumption.

(The main disadvantages for both disruptive and non-disruptive Partial Reconfiguration

techniques are the lack of tools and design software available, or when available, the complexity of the tools.)

- **Self-Reconfiguration:** it is referred as a separate technique, although it is a partial reconfiguration one. In this technique, the FPGA will control its own reconfiguration process and have its own decisions. It is a more complex type of partial reconfiguration, usually augmented by artificial intelligence algorithms [7].

In the figure above (Figure 1) [4], we can see a simple concept on how the Partial Reconfiguration is working. On the FPGA mapping we can see two types of regions:

- Static region (transparent white)
- Dynamic region (transparent red)

The static region is the part of the FPGA where the bitstream will not change the whole program life. On the other hand, the red portions are called dynamic regions. These are the regions where the partial reconfiguration takes place. We can see for the Region 1 that several bitstreams can be loaded (Bmode1, Bmode2, Bmode3) and for the Region 2 Amode1 and Amode2 bitstreams can be loaded, resulting in total of six possible combinations for the configuration [4].

## 2.2. Granularity

When talking about partial reconfiguration, an important property of the system is the granularity [1]. This is defined as the smallest function that can be replaced (modified) in the design. This can be classified into three types:

- Coarse granularity: complete IP cores are modified (reconfigured)
- Medium granularity: only parts of IP cores are reconfigured (register level)
- Low granularity: single or several bits are modified

Low granularity offers the system high versatility, but at the same time it will increase the reconfiguration time. On the other hand, coarse granularity will speed-up the reconfiguration time, but the system will lower its flexibility.

## 2.3. System coupling types

Usually, a reconfigurable system consists in a microprocessor and a reconfigurable device (FPGA reconfigurable material). The microprocessor must run the basic tasks, peripherals control, or the operating system, while the FPGA will take care of the high resource demanding tasks, this way an optimum parallelism is achieved. In Figure 2 we can see several coupling schemes [3].

A more detailed classification of the coupling schemes can be seen below [3]:

- **External stand-alone:** the reconfigurable hardware is connected to the microprocessor inputs and outputs, and it is fully independent. The microprocessor is the host, and the FPGA is the functional unit tied to the main microprocessor data receiver.
- **Attached processing unit:** in this case, the microprocessor is the host, and the FPGA is a coprocessor unit. Both of them are independent and the FPGA is used for leveraging

heavy computing tasks off the microprocessor. The FPGA does not need the supervision of the microprocessor.

- **Coprocessor:** between the first two types. The reconfigurable fabric is located in the data cache, and it can be used either as cache resource or as computing resource.
- **Reconfigurable functional unit:** it permits to define custom instructions as a part of the microprocessor unit.
- **Microprocessor embedded in the FPGA fabric:** the microprocessor can be a hardware core that is burned definitely in the FPGA (cannot be reconfigured) or either a soft processor core that can be replaced by reconfiguring that region before run-time (static region).

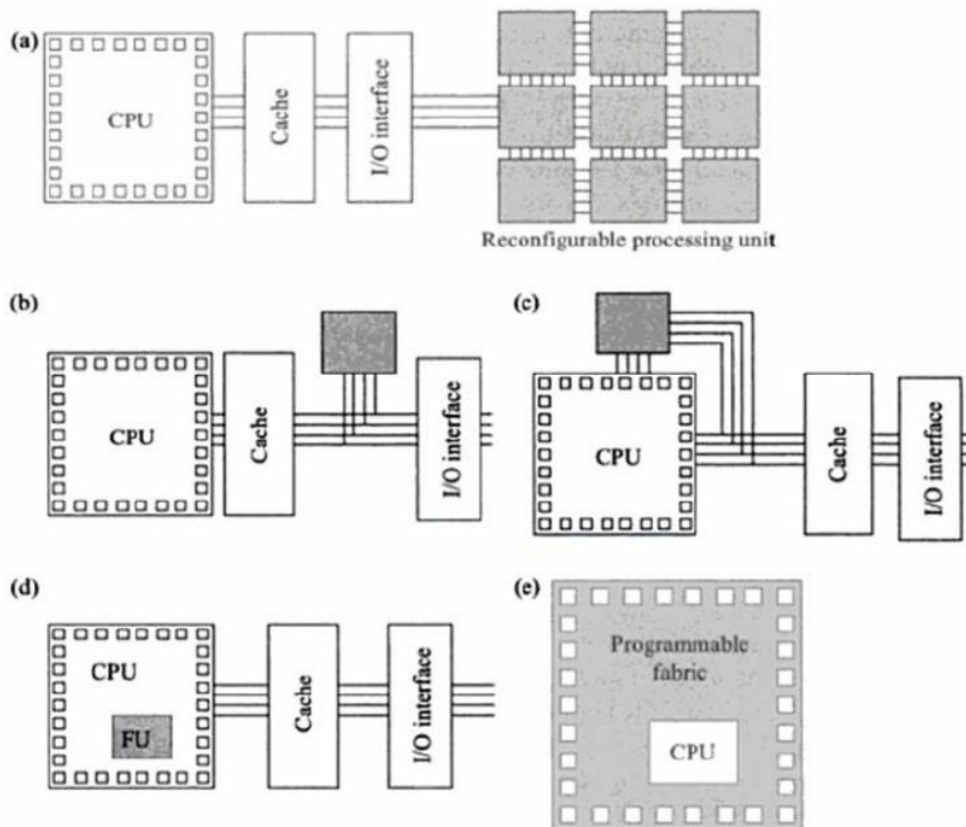


Fig. 2. Coupling Schemes [3]

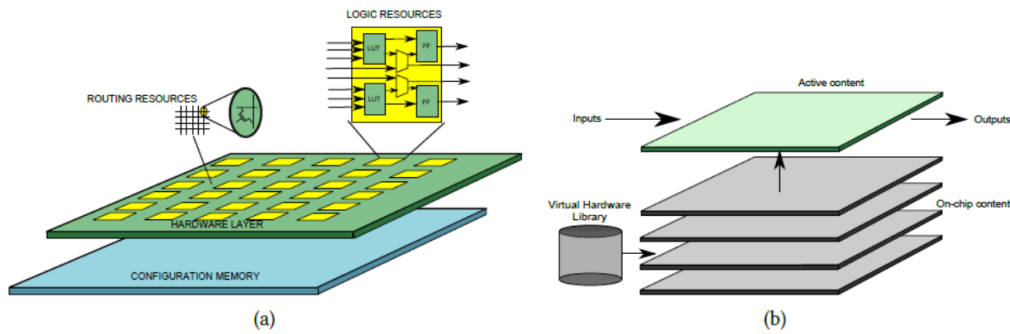
## 2.4. Reconfiguration rate of system

This parameter defines how often the system must be reconfigured [1]. There are three main reconfiguration rate classes [7]:

- **Context-switch:** reconfiguration occurs in order to change some parts of the system, or to change the context of the operating system. This is done at a very low rate, meaning once a day, month, or even year. Because of this, the reconfiguration delays are not important.
- **Runtime:** the reconfiguration is done at the beginning of a task while the other tasks are still running in parallel (multi-context) or even when the system does not have any tasks defined (single-context). It requires a low reconfiguration rate, and it must not affect the system or other running tasks integrity.
- **Realtime:** at the runtime, the reconfiguration time must be precise, and the process must be transparent. Also, the reconfiguration must be masked (all other tasks dependent on the reconfigurable fabric must not be running). For the moment, this type is not strongly supported, but its complexity could make it very precise in real-time applications.

### 3. Architectures

Effective logic capacity increasing and reconfiguration time reducing are the main desired topics when talking about Dynamic Partial Reconfiguration. In the early years, the major constraints of implementing large designs were the limited resource available in the FPGA and the slow reconfiguration caused by fetching big configuration files [2] [9].



**Fig. 3.** a) Typical FPGA architecture. b) Multi-context FPGA architecture with several configuration memory layers [2]

In the above Figure 3 b) we see how the early dynamic reconfiguration architectures dealt with these limitations by increasing the number of configuration context layers. The last generation FPGAs does not lack of hardware resources. In the early days, the partial reconfiguration was used due to lack of resources, but now its target is, for example, to keep alive communications during system reconfiguration, or even sharing the reconfigurable resources between multiple users.

#### 3.1. Academic and non-commercial architectures

The first dynamic reconfiguration system was filed by Xilinx in 1995 as a patent for an FPGA that can store several configurations at the same time. In the first design, two configuration memory arrays were available in the FPGA, and they could store different configuration

bitstreams. This idea was enhanced in 1997 with a time multiplexed FPGA based on the Xilinx XC4000E product.


### GARP architecture

It was a dynamic reconfiguration architecture based on a standard MIPS processor combined with reconfigurable hardware [2].

The reconfigurable hardware was used as a slave, and it was placed on the same configuration layer with the microprocessor, as seen in Figure 4 a.

The program running on the processor oversaw loading and executing the reconfiguration from the configuration array. Also, the reconfigurable fabric had access to the standard memory hierarchy of the processor. Each reconfigurable array was divided into blocks. One block was used for control and the others for logic operations.

The architecture used partial array configuration as individual blocks.



**Fig. 4.** a) GARP architecture. b) Xilinx XC6200 architecture [2]

## 3.2. Commercial Devices supporting Partial Reconfiguration

In the present, only Xilinx and Altera are supporting Partial Reconfiguration architectures [2]. We will only focus on the Xilinx FPGAs.

### Xilinx FPGAs

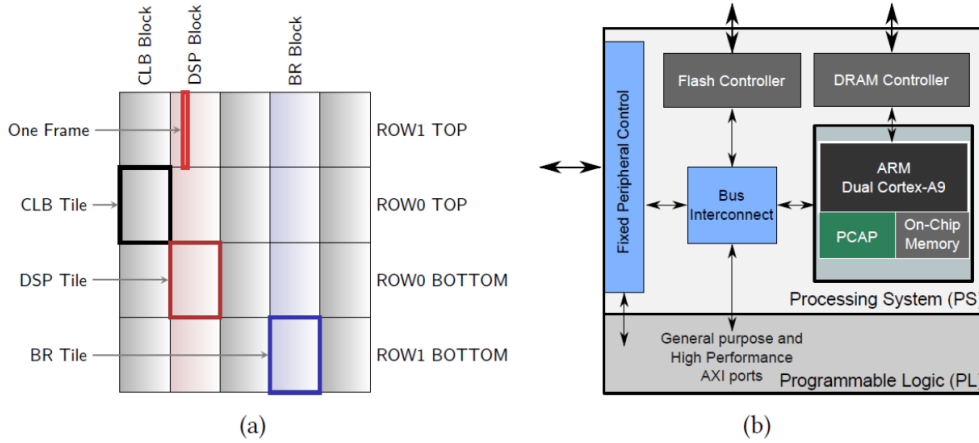
Xilinx had supported Partial Reconfiguration for two decades, so they are the most popular manufacturers when talking about this kind of applications [7].

The first Partial Reconfiguration FPGA developed by Xilinx was the XC6200 series seen in Figure 4 b. This contained only one configuration memory layer and the architecture was divided into functional cells found inside divided tiles-like memory regions. The functional cells contained multiplexers for combinational logic, and routing resources.

The Virtex-II and Virtex-II Pro family made the Partial Reconfiguration more popular. In this architecture, the reconfigurable fabric was arranged in column fashion. These FPGA column primitives contains configurable logic clocks (CLBs), Block RAMs (BRAMs) and multipliers. The CLBs are made of look up-tables (LUTs) and Flip-Flops in two slices [5].

ICAP (Internal Configuration access Port) has been introduced by the Virtex devices family. The ICAP has been a major step, because it allows the FPGA fabric to load partial bitstreams. The bitstreams can be fetched from outside of the FPGA fabric with the help of a soft-processor or a custom state machine that is already loaded and running onto the FPGA fabric. This will result in a circuit implementation that can modify itself [5].

The latest Xilinx FPGAs supporting Dynamic Partial Reconfiguration are the Virtex-5 and Zynq-7000 device families. Both architectures can be seen in Figure 5 [5].



**Fig. 5.** a) Xilinx Virtex architecture. b) Zynq SoC architecture [5]

#### 4. Development tools for Partial Reconfiguration

The main FPGA manufacturers that support Partial Reconfiguration are Xilinx and Altera. They also offer support for the development of Partial Reconfiguration synthetization. In the table below [2] (Figure 6) we can see the main features of Xilinx, Altera and other vendors for the software that they are offering for Partial Reconfiguration [2] [6].

A white dot means that operation is fully manual or not supported, a black and white dot means partial automation or support provided by the tool and a full black dot means that the step is fully automated by the tool requiring no designer intervention.



Tool	<i>High-level Spec.</i>	<i>Partitioning</i>	<i>Floorplanning</i>	<i>Physical implementation</i>	<i>Circuit relocation</i>	<i>Run-time Mgmt.</i>
Xilinx PlanAhead	○	○	◐	●	○	○
Xilinx Vivado	○	○	◐	●	◐	○
Altera Quartus	○	○	◐	●	○	○
Xilinx SDAccel	●	○	N/A	●	N/A	●
Altera OpenCL	●	○	N/A	●	N/A	●
OpenPR	○	○	○	◐	●	○
CoPR	◐	●	◐	○	○	◐
GoAhead	○	○	◐	◐	●	○
Caronte	◐	◐	○	○	◐	○
GePaRD	●	○	○	○	○	○
PaRAT	●	●	○	○	○	○
OSSS+R	●	◐	○	○	○	◐

**Fig. 6.** Comparison features supported by different PR tools [2]

## 5. Conclusion and further work

In the last years, the ASIC world has been surpassed by the FPGAs in the custom hardware and low mass production field. The need of fast developing computing intensive embedded devices has led to a progress increase in the FPGA research and development.

The Dynamic (run-time) Partial Reconfiguration on FPGAs topic was opened in 1995 because of the lack of hardware resource on the reconfigurable fabric, but these days this topic is growing fueled by other necessities, like the need to share the same resource with multiple users, switch hungry tasks by rearranging parts of the hardware subsystems in multitasking, save power by disabling subsystems, creating self-recovering systems that can rebuild its own subsystems or even for fault tolerant systems that needs extra protection.

In conclusion, the Dynamic Partial Reconfiguration on FPGAs is a promising field that can accomplish hungry computing tasks withing a low power consumption that need to be further studied.

## References

- [1] Y.E. Krasteva. Reconfigurable Computing Based on Commercial FPGAs. Solutions for the Design and Implementation of Partially Reconfigurable Systems, 2009
- [2] KIZHEPPATT VIPIN and SUHAIB A. FAHMY. 2017. FPGA Dynamic and Partial Reconfiguration: A Survey of Architectures, Methods, and Applications. ACM Comput. Surv. 1, 1, Article 1 (January 2017)
- [3] Bashir Al-Hashimi. System-on-chip: Next Generation Electronics. IEE, 2006
- [4] KIZHEPPATT VIPIN, Partial reconfiguration (youtube tutorial series) 2020
- [5] Vivado Design Suite User Guide Partial Reconfiguration - UG909 (v2018.1), 2018
- [6] Fernando Moraes, DYNAMIC AND PARTIAL RECONFIGURATION IN FPGA SOCS: REQUIREMENTS TOOLS AND A CASE STUDY
- [7] Oswald Berthold, Self-reconfiguring System-on-Chip using Linux on a Virtex-5 FPGA, 2012
- [8] Rashad S. Oreifej, A Sustainable Autonomic Architecture for Organically Reconfigurable Computing Systems, 2011
- [9] Ahmed Al-Wattar, Efficient Scheduling, Mapping and Resource Prediction for Dynamic Run time Operating Systems, 2015